

Edge-Weighted Hypergraph Transversals & Contextuality

Thomas C. Fraser^{1,2,*}

¹Perimeter Institute for Theoretical Physics, Waterloo, Ontario, Canada, N2L 2Y5

²University of Waterloo, Waterloo, Ontario, Canada, N2L 3G1

(Dated: September 13, 2017)

This is the abstract.

CONTENTS

I. Introduction	1
A. Applications	1
II. Marginal Satisfiability	1
A. Definitions	1
B. Linearity	2
C. Marginal Polytopes	3
D. Logical Contextuality	4
III. An Observation	4
A. An Antecedent Hierarchy	4
B. The Antecedent Hypergraph	4
C. Irreducibility	5
D. Marginal Symmetries	5
E. Curated Inequalities	5
F. Targeted Searches	5
G. Relaxations	5
IV. Edge-Weighted Hypergraph Transversals	5
A. Preliminaries	5
B. Hypergraph Transversals	5
C. Adding Weights	5
D. Unsolved	5
V. Conclusions	5
Acknowledgments	5
References	5

I. INTRODUCTION

A. Applications

II. MARGINAL SATISFIABILITY

A. Definitions

To every random variable¹ v there corresponds a prescribed set of **outcomes** \mathcal{O}_v and a set of **events over** v

denoted $\Omega(v)$ corresponding to the set of all functions of the form $\omega : \{v\} \rightarrow \mathcal{O}_v$. Evidently, $\Omega(v)$ and \mathcal{O}_v are isomorphic structures and their distinction can be confounding. There is rarely any harm in referring synonymously to either as outcomes. Nonetheless, a sheaf-theoretic treatment of contextuality [1] demands the distinction. Specifically for this work, the distinction becomes essential for the exploitation of marginal symmetries in Section III D. As a natural generalization we define the event over a collection of random variables $V = \{v_1, \dots, v_n\}$ in a parallel manner:

$$\Omega(V) \equiv \{\omega : V \rightarrow \mathcal{O}_V \mid \forall v \in V, \omega(v) \in \mathcal{O}_v\}$$

Furthermore, the **domain** $\mathcal{D}(\omega)$ of an event ω is the set of random variables it valuates, i.e. if $\omega \in \Omega(V)$ then $\mathcal{D}(\omega) = V$.

For every $V' \subset V$ and $\omega \in \Omega(V)$, the **restriction of ω onto V'** (denoted $\omega|_{V'}$) corresponds to the unique event in $\Omega(V')$ that agrees with ω for all valuations of variables in V' , i.e. $\forall v' \in V' : \omega|_{V'}(v') = \omega(v')$. Using this notational framework, a probability distribution or simply **distribution** \mathbf{p}_V is a probability measure on $\Omega(V)$, assigning to each $\omega \in \Omega(V)$ a real number $\mathbf{p}_V(\omega) \in [0, 1]$ such that $\sum_{\omega \in \Omega(V)} \mathbf{p}_V(\omega) = 1$. The set of all distributions over $\Omega(V)$ is denoted \mathcal{P}_V . Moreover, given $\mathbf{p}_V \in \mathcal{P}_V$ and $V' \subset V$, there is an induced distribution $\mathbf{p}_V|_{V'} \in \mathcal{P}_{V'}$ obtained by *marginalizing* \mathbf{p}_V :

$$\mathbf{p}_V|_{V'}(\omega') = \sum_{\substack{\omega \in \Omega(V) \\ \omega|_{V'} = \omega'}} \mathbf{p}_V(\omega) \quad (1)$$

Presently, the reader is equipped with sufficient notation and terminology to comprehend the marginal problem.

Definition 1. The Marginal Problem: Given a collection of m distributions $\{\mathbf{p}_{V_1}, \dots, \mathbf{p}_{V_m}\}$, does there exist a distribution $\mathbf{p}_\Lambda \in \mathcal{P}_\Lambda$ with $\Lambda \equiv \bigcup_{i=1}^m V_m$ such that $\forall i : \mathbf{p}_\Lambda|_{V_i} = \mathbf{p}_{V_i}$?

To facilitate further discussion of this problem, several pieces of nomenclature will be introduced. First, the set $\mathcal{V} = \{V_1, \dots, V_m\}$ is called the **marginal scenario** while its elements are called the **marginal contexts**. The collection of distributions $\mathbf{p}_{*\mathcal{V}} \equiv \{\mathbf{p}_{V_1}, \dots, \mathbf{p}_{V_m}\}$ ² is

* tcfrazer@tcfrazer.com

¹ Throughout this document, it is assumed that all random variables are discrete and have finite cardinality.

² The subscript $*$ preceding \mathcal{V} is added for clarity; $\mathbf{p}_{*\mathcal{V}}$ is *not* a distribution but a set of distributions over \mathcal{V} . The $_{*\mathcal{V}}$ convention is adopted throughout this report.

called the **marginal model** [2]³. The distribution \mathbf{p}_Λ , if it exists, is termed the **joint distribution**. Strictly speaking, as defined by [2], a marginal scenario forms an *abstract simplicial complex*, meaning it satisfies the supplementary requirement that all subsets of contexts are also contexts, i.e. $\forall V \in \mathcal{V} : V' \subset V \implies V' \in \mathcal{V}$. Throughout this work, we exclusively consider (without loss of generality) *maximal* marginal scenarios, restricting our focus to the contexts which are contained in no others. Finally, a marginal model $\mathbf{p}_{*\mathcal{V}}$ is said to be **non-contextual**, and will be denoted $\mathbf{p}_{*\mathcal{V}} \in \mathcal{N} \subseteq \mathcal{P}_{*\mathcal{V}}$ if it admits a joint distribution and **contextual** otherwise ($\mathbf{p}_{*\mathcal{V}} \in \mathcal{N}$). Equipped with this additional terminology and notation, the marginal problem now reads: given $\mathbf{p}_{*\mathcal{V}}$, is $\mathbf{p}_{*\mathcal{V}} \in \mathcal{N}$ or not?

B. Linearity

An essential feature of the marginal problem is linearity; the marginalization of \mathbf{p}_Λ onto the marginal contexts $\{\mathbf{p}_\Lambda|_V \mid V \in \mathcal{V}\}$ is a linear transformation, requiring only the summations pursuant to Eq. (1). Consequently, it is advantageous to consider the statement of the marginal problem as a matrix multiplication. To this end, for each marginal scenario \mathcal{V} we define a binary matrix \mathcal{M} called the **incidence matrix** which implements this mapping. The columns of \mathcal{M} are indexed by *joint events* $j \in \Omega(\Lambda)$ and the rows are indexed by *marginal events* $m \in \Omega(V)$ for some $V \in \mathcal{V}$. By deliberate abuse of notation, we will denote the set of all marginal events as $\Omega(*\mathcal{V})$ and is defined as the following disjoint union:

$$\Omega(*\mathcal{V}) \equiv \coprod_{V \in \mathcal{V}} \Omega(V)$$

The $|\Omega(*\mathcal{V})| \times |\Omega(\Lambda)|$ incidence matrix \mathcal{M} is then defined element-wise for $m \in \Omega(*\mathcal{V})$ and $j \in \Omega(\Lambda)$:

$$\mathcal{M}_j^m = \begin{cases} 1 & j|_{\mathcal{D}(m)} = m \\ 0 & \text{otherwise} \end{cases} \quad (2)$$

Conceptually, the entries of this matrix are populated with ones whenever the marginal event (row) m is the restriction of some joint event (column) j . For a given marginal scenario \mathcal{V} , \mathcal{M} represents the tuple of restriction maps $\mathcal{M} : \Omega(\Lambda) \rightarrow \prod_{V \in \mathcal{V}} \Omega(V) :: j \mapsto \{j|_V \mid V \in \mathcal{V}\}$ [1]. Furthermore, note that the component indices of \mathcal{M} in Eq. (2) are deliberately separated. Among other reasons, this is done to allow one to denote the m -th row of \mathcal{M} as \mathcal{M}^m and the j -th column as \mathcal{M}_j . For further notational convenience, since \mathcal{M} is a binary matrix, we let \mathcal{M}^m and \mathcal{M}_j analogously correspond their respective *supports*⁴,

e.g. $m \in \sigma(\mathcal{M}_j)$ if and only if $\mathcal{M}_j^m = 1$. Throughout the remainder of this report, the utility of the incidence matrix \mathcal{M} will be indispensable.

To illustrate this concretely, consider the following example. Let Λ be 3 binary variables $\{a, b, c\}$ and \mathcal{V} be the marginal scenario $\mathcal{V} = \{\{a, b\}, \{b, c\}, \{a, c\}\}$. The incidence matrix for \mathcal{V} becomes:

$$\begin{array}{l} (a, b, c) \mapsto \\ (a \rightarrow 0, b \rightarrow 0) \\ (a \rightarrow 0, b \rightarrow 1) \\ (a \rightarrow 1, b \rightarrow 0) \\ (a \rightarrow 1, b \rightarrow 1) \\ (b \rightarrow 0, c \rightarrow 0) \\ (b \rightarrow 0, c \rightarrow 1) \\ (b \rightarrow 1, c \rightarrow 0) \\ (b \rightarrow 1, c \rightarrow 1) \\ (a \rightarrow 0, c \rightarrow 0) \\ (a \rightarrow 0, c \rightarrow 1) \\ (a \rightarrow 1, c \rightarrow 0) \\ (a \rightarrow 1, c \rightarrow 1) \end{array} \mapsto \begin{pmatrix} (0,0,0) & (0,0,1) & (0,1,0) & (0,1,1) & (1,0,0) & (1,0,1) & (1,1,0) & (1,1,1) \\ \mathbf{1} & \mathbf{1} & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & \mathbf{1} & \mathbf{1} & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & \mathbf{1} & \mathbf{1} & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & \mathbf{1} & \mathbf{1} \\ \mathbf{1} & 0 & 0 & 0 & \mathbf{1} & 0 & 0 & 0 \\ 0 & \mathbf{1} & 0 & 0 & 0 & \mathbf{1} & 0 & 0 \\ 0 & 0 & \mathbf{1} & 0 & 0 & 0 & \mathbf{1} & 0 \\ 0 & 0 & 0 & \mathbf{1} & 0 & 0 & 0 & \mathbf{1} \\ \mathbf{1} & 0 & \mathbf{1} & 0 & 0 & 0 & 0 & 0 \\ 0 & \mathbf{1} & 0 & \mathbf{1} & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & \mathbf{1} & 0 & \mathbf{1} & 0 \\ 0 & 0 & 0 & 0 & 0 & \mathbf{1} & 0 & \mathbf{1} \end{pmatrix} \quad (3)$$

In addition, for any joint distribution $\mathbf{p}_\Lambda \in \mathcal{P}_\Lambda$ we associate a joint distribution *vector* \mathbf{p}_Λ (identically denoted) indexed by $j \in \Omega(\Lambda)$, i.e. $\mathbf{p}_\Lambda^j \equiv \mathbf{p}_\Lambda(j)$. Analogously, for each marginal model $\mathbf{p}_{*\mathcal{V}} \in \mathcal{P}_{*\mathcal{V}}$ there is an associated marginal distribution *vector* $\mathbf{p}_{*\mathcal{V}}$ indexed by $m \in \Omega(*\mathcal{V})$ such that $\mathbf{p}_{*\mathcal{V}}^m \equiv \mathbf{p}_{\mathcal{D}(m)}(m)$. Using these vectors, the marginal problem becomes the following linear program:

Definition 2. The Marginal Linear Program (MLP):

$$\begin{aligned} & \text{minimize: } \emptyset \cdot \mathbf{p}_\Lambda^5 \\ & \text{subject to: } \mathbf{p}_\Lambda \geq 0 \\ & \quad \mathcal{M} \cdot \mathbf{p}_\Lambda = \mathbf{p}_{*\mathcal{V}} \end{aligned} \quad (4)$$

As such, $\mathbf{p}_{*\mathcal{V}} \in \mathcal{N}$ if and only if MLP is a *feasible* linear program. Importantly, if MLP is feasible, it will return the joint distribution \mathbf{p}_Λ . To every linear program, there exists a dual linear program that characterizes the feasibility of the original [3]. Constructing the dual linear program is a well-defined procedure [4].

Definition 3. The Dual Marginal Linear Program (DMLP):

$$\begin{aligned} & \text{minimize: } \gamma \cdot \mathbf{p}_{*\mathcal{V}} \\ & \text{subject to: } \gamma \cdot \mathcal{M} \geq 0 \end{aligned}$$

By construction, DMLP completely determines the whether or not $\mathbf{p}_{*\mathcal{V}} \in \mathcal{N}$ or not. If $\mathbf{p}_{*\mathcal{V}} \in \mathcal{N}$, then MLP is feasible and the following holds,

$$\gamma \cdot \mathbf{p}_{*\mathcal{V}} = \gamma \cdot (\mathcal{M} \cdot \mathbf{p}_\Lambda) \geq 0 \quad (5)$$

³ In [1], $\mathbf{p}_{*\mathcal{V}}$ is instead called an *empirical model*.

⁴ The *support* $\sigma(f)$ of a mapping f is the subset of its domain $\mathcal{D}(f)$ that is not mapped to a zero element: $\sigma(f) = \{x \in \mathcal{D}(f) \mid f(x) \neq 0\}$.

⁵ Note that the primal value of the linear program is of no interest, all that matters is its *feasibility*. Here \emptyset denotes a null vector of all zero entries.

because both $\gamma \cdot \mathcal{M} \geq 0$ and $\mathbf{p}_\Lambda \succeq 0$. If however, $\gamma \cdot \mathbf{p}_{*\mathcal{V}} < 0$, then Eq. (5) is violated and $\mathbf{p}_{*\mathcal{V}} \notin \mathcal{N}$ ⁶. In summary, the sign of $d \equiv \min(\gamma \cdot \mathbf{p}_{*\mathcal{V}})$ answers the marginal program; $\mathbf{p}_{*\mathcal{V}} \in \mathcal{N}$ if and only if $d \geq 0$ ⁷.

Corollary 1. *All linear, homogeneous constraints $\gamma \cdot \mathbf{p}_{*\mathcal{V}} \geq 0$ constraining non-contextual marginal models $\mathcal{N} \subseteq \mathcal{P}_{*\mathcal{V}}$ satisfy $\gamma \cdot \mathcal{M} \geq 0$. Moreover, all vectors γ satisfying $\gamma \cdot \mathcal{M} \geq 0$ correspond to valid constraints $\gamma \cdot \mathbf{p}_{*\mathcal{V}} \geq 0$ for $\mathcal{N} \subseteq \mathcal{P}_{*\mathcal{V}}$.*

In light of Definitions 2 and 3, when supplied with a particular marginal model $\mathbf{p}_{*\mathcal{V}}$, the marginal problem can be solved computationally by evaluating DMLP to determine the feasibility of MLP. A more difficult variant of the marginal problem is one wherein no particular marginal model is supplied.

Definition 4. The General Marginal Problem (GMP): Given a marginal scenario \mathcal{V} find a set of independent constraints Γ which completely constraint $\mathcal{N} \subseteq \mathcal{P}_{*\mathcal{V}}$; i.e. $\mathbf{p}_{*\mathcal{V}} \in \mathcal{N}$ if and only if it satisfies all constraints in Γ .

The remainder of this paper is concerned with methods for solving (or partially solving) GMP. Specifically, Section II C discusses existing methods for completely solving GMP and outlines some of their disadvantages. Section II D summarizes an existing method for completely solving a possibilistic variant of GMP. Sections II C and II D motivate Section III, wherein a new method for completely solving GMP is presented.

C. Marginal Polytopes

The complete space of marginal models over \mathcal{V} (denoted $\mathcal{P}_{*\mathcal{V}}$) can be partitioned into two spaces: the contextual marginal models ($\bar{\mathcal{N}}$) and the non-contextual marginal models ($\mathcal{N} \equiv \mathcal{P}_{*\mathcal{V}} \setminus \bar{\mathcal{N}}$). Pitowsky [7] demonstrates that \mathcal{N} forms a *convex* polytope commonly referred to as the **marginal polytope** for \mathcal{V} . When embedded in $\mathbb{R}^{|\Omega(\mathcal{V})|}$, the extremal rays of the marginal polytope correspond to the columns of \mathcal{M} which further correspond to all *deterministic* joint distributions $\mathbf{p}_\Lambda \in \mathcal{P}_\Lambda$ ⁸. The normalization of \mathbf{p}_Λ ($\sum_j \mathbf{p}_\Lambda^j = 1$) defines the convexity of the polytope; each marginal model $\mathbf{p}_{*\mathcal{V}} \in \mathcal{P}_{*\mathcal{V}}$ must be a convex mixture of the deterministic marginal models pursuant to Eq. (4). Consequently,

characterizing the contextuality of marginal models is manifestly a problem of polytope description. Notably, the **facets** of a marginal polytope correspond to a finite set of linear inequalities that are complete in the sense that all contextual distributions violate at least one facet inequality [8]. From the perspective of a marginal polytope, convex hull algorithms or linear quantifier elimination can be used to compute a representation of the complete set of facet inequalities and consequently completely solve the GMP. A popular tool for linear quantifier elimination is *Fourier-Motzkin elimination* [9–12]. In this report, we will avoid expounding upon the Fourier-Motzkin procedure and instead recall a few of its notable features and consequences.⁹

Definition 5. [3, Section 12.2] Given a system of linear inequality constraints $\mathcal{S} = \{A \cdot x \leq b\}$ constraining some free variables x , the **Fourier-Motzkin elimination** procedure eliminates some of the variables in x and returns a system of linear inequality constraints $\mathcal{S}' = \{A' \cdot x' \leq b'\}$ over $x' \subset x$ such that any solution x' of \mathcal{S}' will permit at least one compatible solution x of \mathcal{S} (and vice versa).

$$\exists x' : A' \cdot x' \leq b' \iff \exists x : A \cdot x \leq b \quad (6)$$

In particular, the following system of linear inequalities defines the marginal polytope for \mathcal{V} :

$$\begin{aligned} \forall m \in \Omega(\mathcal{V}) : \quad & \mathbf{p}_{*\mathcal{V}}^m - \sum_j \mathcal{M}_j^m \mathbf{p}_\Lambda^j \geq 0 \\ \forall m \in \Omega(\mathcal{V}) : \quad & -\mathbf{p}_{*\mathcal{V}}^m + \sum_j \mathcal{M}_j^m \mathbf{p}_\Lambda^j \geq 0 \\ \forall j \in \Omega(\Lambda) : \quad & \mathbf{p}_\Lambda^j \geq 0 \\ & \sum_j \mathbf{p}_\Lambda^j \geq 1 \\ & -\sum_j \mathbf{p}_\Lambda^j \geq -1 \end{aligned} \quad (7)$$

Using the Fourier-Motzkin elimination procedure, it is possible to eliminate the variables \mathbf{p}_Λ^j relating to joint events and obtain a system of linear inequalities constraining only marginal events $\mathbf{p}_{*\mathcal{V}}^m$ which completely characterizes the set of non-contextual marginal models $\mathcal{N} \subseteq \mathcal{P}_{*\mathcal{V}}$.

Lemma 1.¹⁰ *There exists a finite set of integral vectors Γ such that for all $\mathbf{p}_{*\mathcal{V}} \in \mathcal{P}_{*\mathcal{V}}$:*

$$\mathbf{p}_{*\mathcal{V}} \in \mathcal{N} \iff \forall \gamma \in \Gamma : \gamma \cdot \mathbf{p}_{*\mathcal{V}} \geq 0 \quad (8)$$

Proof. The finiteness and existence of Γ is a fundamental property of polytopes [9, 11–14]. The fact that each vector $\gamma \in \Gamma$ need only be integer-valued follows from

⁶ These observations are collectively a consequence of Farkas's Lemma [5].

⁷ In particular, if $d \geq 0$, then $d = 0$ due to the existence of the trivial solution $\gamma = \emptyset$. This observation is an instance of the *Complementary Slackness Property* [6]. Alternatively, if $d < 0$, then it is unbounded $d = -\infty$ due to the *Unbounded Property* [6].

⁸ A deterministic distribution \mathbf{p}_Λ is a distribution in which a singular event $j \in \Omega(\Lambda)$ occurs with certainty, i.e. $\mathbf{p}_\Lambda^j = 1$ and $\forall j' \neq j : \mathbf{p}_\Lambda^{j'} = 0$.

⁹ Applying the Fourier-Motzkin procedure to completely solve GMP is discussed in more detail in Fritz and Chaves [2].

¹⁰ This is a stronger variant of [11, Proposition 7].

the integer-valued coefficients that constrain Eq. (7). Finally, the homogeneity of the constraints in Eq. (8) follows from the assumption that each $\mathbf{p}_{*\mathcal{V}} \in \mathcal{P}_{*\mathcal{V}}$ a priori satisfies normalization constraints context-wise; i.e. $\forall V \in \mathcal{V} : \sum_{m \in \Omega(V)} \mathbf{p}_V^m = 1$ ¹¹. \square

D. Logical Contextuality

Let $a \in \Omega_*(\mathcal{V})$ be *any* marginal event and $C = \{c_1, \dots, c_n\} \subseteq \Omega_*(\mathcal{V})$ be a subset of marginal events such that the following logical implication holds for *all* marginal models $\mathbf{p}_{*\mathcal{V}} \in \mathcal{P}_{*\mathcal{V}}$:

$$a \implies c_1 \vee \dots \vee c_n = \bigvee_{c \in C} c \quad (9)$$

Which can be dictated: *whenever the event a occurs, at least one event in C occurs.* In accordance with the logical form of Eq. (9), a will be referred to as the **antecedent** and C as the **consequent set**. To clarify, a marginal model $\mathbf{p}_{*\mathcal{V}} \in \mathcal{P}_{*\mathcal{V}}$ satisfies Eq. (9) if there always at least one $c \in C$ that is *possible* ($\mathbf{p}_{*\mathcal{V}}^c > 0$) whenever a is possible. A marginal model violates Eq. (9) whenever *none* of events in C are possible while a remains possible. Marginal models that violate logical statements such as Eq. (9) are known as **Hardy Paradoxes** [10, 15, 16]. Motivated by a greater sense of robustness compared to possibilistic constraints, the concept of witnessing quantum contextuality on a logical level has been analyzed thoroughly for decades [11, 17].

All logical implications of the form of Eq. (9) can be derived by first selecting an antecedent marginal event a , then constructing a consequent set C such that Eq. (9) holds. This is accomplished by making use of Lemma 2.

Lemma 2. *Let $m \in \Omega_*(\mathcal{V})$ be a marginal event. Then for all non-contextual marginal models $\mathbf{p}_{*\mathcal{V}} \in \mathcal{N}$,*

$$m \iff \bigvee_{j \in \sigma(\mathcal{M}^m)} j$$

Essentially, if a joint distribution does exist ($\mathbf{p}_{*\mathcal{V}} \in \mathcal{N}$), then event m represents partial knowledge of the entire system of variables Λ ; whenever m occurs, exactly one joint event j has occurred in reality and m must be a restriction of j . Applying Lemma 2 to the antecedent $a \in \Omega_*(\mathcal{V})$ and consequent set $C \subseteq \Omega_*(\mathcal{V})$,

$$\begin{aligned} a &\iff \bigvee_{j \in \sigma(\mathcal{M}^a)} j \\ \bigvee_{c \in C} c &\iff \bigvee_{c \in C} \bigvee_{j \in \sigma(\mathcal{M}^c)} j \end{aligned} \quad (10)$$

Therefore, if a subset C of $\Omega_*(\mathcal{V})$ (preferably excluding a) can be found such that,

$$\sigma(\mathcal{M}^a) \subseteq \bigcup_{c \in C} \sigma(\mathcal{M}^c) \quad (11)$$

then Eq. (9) follows from Eqs. (10,11).

It is possible to show that for each logical constraint of the form Eq. (9), there exists a corresponding probabilistic constraint that is tighter. Corollary 2¹² generalizes Lemma 2.

Corollary 2. *Let $m \in \Omega_*(\mathcal{V})$ be a marginal event. Then for all non-contextual marginal models $\mathbf{p}_{*\mathcal{V}} \in \mathcal{N}$,*

$$\mathbf{p}_{*\mathcal{V}}^m = \sum_{j \in \sigma(\mathcal{M}^m)} \mathbf{p}_\Lambda^j$$

III. AN OBSERVATION

A. An Antecedent Hierarchy

B. The Antecedent Hypergraph

Given an antecedent multi-set γ where $\gamma \preceq 0$, we identify the **inhibiting set** of joint events $\mathcal{I}(\gamma) \subseteq \Omega(\Lambda)$ preventing $\gamma \cdot \mathcal{M}$ from being positive semi-definite:

$$\mathcal{I}(\gamma) \equiv \left\{ j \in \Omega(\Lambda) \mid (\gamma \cdot \mathcal{M})_j < 0 \right\}$$

The inhibiting set $\mathcal{I}(\gamma)$ of γ completely characterizes the **antecedent hypergraph** $\mathcal{H}(\gamma)$ whose edges \mathcal{E}_j are indexed by the inhibiting events $j \in \mathcal{I}(\gamma)$. Each edge $\mathcal{E}_j \subseteq \Omega_*(\mathcal{V})$ corresponds to the set of the marginal events $m \in \Omega_*(\mathcal{V})$ which are restrictions of j . Specifically,

$$\begin{aligned} \mathcal{H}(\gamma) &\equiv \{ \mathcal{E}_j \mid j \in \mathcal{I}(\gamma) \} \\ \mathcal{E}_j &\equiv \{ m \in \Omega_*(\mathcal{V}) \mid m = j|_{\mathcal{D}(m)}, \gamma_m = 0 \} \end{aligned}$$

Todo (TC Fraser): Incorporate weights

$$\omega(\gamma) = \left\{ \omega_j = -(\gamma \cdot \mathcal{M})_j \mid j \in \mathcal{I}(\gamma) \right\}$$

Todo (TC Fraser): Define set of right-minimal inequalities

$$\mathcal{I}(\gamma) = \left\{ (\gamma + \tilde{\gamma}) \cdot \mathbf{p}_{*\mathcal{V}} \geq 0 \mid \tilde{\gamma} \in \text{Tr}_{\omega(\gamma)}[\mathcal{H}(\gamma)] \right\}$$

¹¹ Specifically, any inhomogeneous constraint $\gamma \cdot \mathbf{p}_{*\mathcal{V}} \geq \alpha$ can be *homogenized* by replacing α with $\sum_{m \in \Omega(V)} \alpha \mathbf{p}_V^m$.

¹² Corollary 2 is simply the m -th row of Eq. (4).




FIG. 1. Dual-representations of a hypergraph $\mathcal{H} = \{\mathcal{E}_1, \mathcal{E}_2, \mathcal{E}_3, \mathcal{E}_4, \mathcal{E}_5\}$.

B. Hypergraph Transversals

C. Adding Weights

D. Unsolved

Consider γ_1, γ_2 and $\gamma_1 + \gamma_2$.

$$I(\gamma_a) = \{(\gamma_a + \tilde{\gamma}_a) \cdot \mathbf{p}_* \mathbf{v} \geq 0 \mid \tilde{\gamma}_a \in \text{Tr}_{\omega(\gamma_a)}[\mathcal{H}(\gamma_a)]\} \quad a = 1, 2$$

$$I(\gamma_1 + \gamma_2) = \{(\gamma_1 + \gamma_2 + \tilde{\gamma}_{12}) \cdot \mathbf{p}_* \mathbf{v} \geq 0 \mid \tilde{\gamma}_{12} \in \text{Tr}_{\omega(\gamma_1) + \omega(\gamma_2)}[\mathcal{H}(\gamma_1 + \gamma_2)]\}$$

C. Irreducibility

D. Marginal Symmetries

E. Curated Inequalities

F. Targeted Searches

G. Relaxations

IV. EDGE-WEIGHTED HYPERGRAPH TRANSVERSALS

A. Preliminaries

V. CONCLUSIONS

ACKNOWLEDGMENTS

-
- [1] S. Abramsky and A. Brandenburger, “The Sheaf-Theoretic Structure Of Non-Locality and Contextuality,” *New J. Phys* **13**, 113036 (2011).
 - [2] T. Fritz and R. Chaves, “Entropic Inequalities and Marginal Problems,” *IEEE Trans. Info. Theor.* **59**, 803 (2011).
 - [3] A. Schrijver, *Theory of Linear and Integer Programming* (Wiley, 1998).
 - [4] S. Lahaie, “How to take the Dual of a Linear Program,” (2008).
 - [5] S. Fang and S. Puthenpura, *Linear Optimization and Extensions: Theory and Algorithms* (Englewood Cliffs, NJ: Prentice-Hall, 1993) p. 60.
 - [6] Bradley, Hax, and Magnanti, *Applied Mathematical Programming* (Addison-Wesley, 1977) Chap. 4, pp. 143–144.
 - [7] I. Pitowsky, “Correlation polytopes: Their geometry and complexity,” *Math. Prog.* **50**, 395 (1991).

- [8] N. Brunner, D. Cavalcanti, S. Pironio, V. Scarani, and S. Wehner, “Bell nonlocality,” *Rev. Mod. Phys.* **86**, 419 (2013).
- [9] G. B. Dantzig and B. C. Eaves, “Fourier-Motzkin elimination and its dual,” *J. Combin. Theor. A* **14**, 288 (1973).
- [10] E. Wolfe, R. W. Spekkens, and T. Fritz, “The Inflation Technique for Causal Inference with Latent Variables,” (2016), [arXiv:1609.00672](https://arxiv.org/abs/1609.00672).
- [11] S. Abramsky and L. Hardy, “Logical Bell inequalities,” *Phys. Rev. A* **85**, 062114 (2012).
- [12] C. Jones, E. C. Kerrigan, and J. Maciejowski, *Equality set projection: A new algorithm for the projection of polytopes in halfspace representation*, Tech. Rep. (Cambridge University Engineering Dept, 2004).
- [13] B. G. Jiri Matousek, *Understanding and Using Linear Programming (Universitext)* (Springer, 2013).
- [14] G. M. Ziegler, *Lectures on Polytopes* (Springer New York, 1995).
- [15] S. Mansfield and T. Fritz, “Hardy’s Non-locality Paradox and Possibilistic Conditions for Non-locality,” *Found. Phys.* **42**, 709 (2011).
- [16] L. Mančinska and S. Wehner, “A unified view on Hardy’s paradox and the Clauser–Horne–Shimony–Holt inequality,” *J. Phys. A* **47**, 424027 (2014).
- [17] D. M. Greenberger, “Bell’s theorem without inequalities,” *Am. J. Phys.* **58**, 1131 (1990).